

CAT_4

May 31, 2023

1 CAT: 4 - CUDA

1.1 1 CUDA Execution Environment

```
[ ]: ! nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2022 NVIDIA Corporation
Built on Wed_Sep_21_10:33:58_PDT_2022
Cuda compilation tools, release 11.8, V11.8.89
Build cuda_11.8.r11.8/compiler.31833905_0
```

```
[ ]: ! pip install \
git+https://github.com/andreinechaev/nvcc4jupyter.git
%load_ext nvcc_plugin
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-
build-tx1m4v7m
  Running command git clone --filter=blob:none --quiet
https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-tx1m4v7m
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit
aac710a35f52bb78ab34d2e52517237941399eff
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py) ... done
  Created wheel for NVCCPlugin: filename=NVCCPlugin-0.0.2-py3-none-any.whl
size=4287
sha256=321c937882579f848dcb06c51984e84778d1374898ca5878cebf5e4d471a50c1
  Stored in directory: /tmp/pip-ephem-wheel-cache-
k7yvvhvrt/wheels/a8/b9/18/23f8ef71ceb0f63297dd1903aedd067e6243a68ea756d6fee
Successfully built NVCCPlugin
Installing collected packages: NVCCPlugin
Successfully installed NVCCPlugin-0.0.2
created output directory at /content/src
Out bin /content/result.out
```

```
[ ]: %%cu

#include <stdio.h>

__global__ void hello_kernel(void) {
    printf("Hello world from the device!\n");
}

int main(void) {
    printf("Hello world from the device!\n");
    hello_kernel<<<1,1>>>();
    cudaDeviceSynchronize();
    return 0;
}
```

```
Hello world from the device!
Hello world from the device!
```

1.2 2 Threading

```
[ ]: %%cu

#include <stdio.h>

__global__ void hello_kernel(void) {
    int blockId = blockIdx.x + blockIdx.y * gridDim.x + blockIdx.z * gridDim.x
    ↪ * gridDim.y;
    int threadId = blockId * blockDim.x + threadIdx.x;
    printf("Hi! My Id is %d, I am the thread %d out of %d in block %d\n",
    ↪ threadIdx, threadIdx.x, blockDim.x, blockId);
}

int main(void) {
    int gridDimension = 6;
    int blockDim = 4;
    hello_kernel<<<gridDimension, blockDim>>>();
    cudaDeviceSynchronize();
    return 0;
}
```

```
Hi! My Id is 8, I am the thread 0 out of 4 in block 2
Hi! My Id is 9, I am the thread 1 out of 4 in block 2
Hi! My Id is 10, I am the thread 2 out of 4 in block 2
Hi! My Id is 11, I am the thread 3 out of 4 in block 2
Hi! My Id is 20, I am the thread 0 out of 4 in block 5
Hi! My Id is 21, I am the thread 1 out of 4 in block 5
Hi! My Id is 22, I am the thread 2 out of 4 in block 5
```

```

Hi! My Id is 23, I am the thread 3 out of 4 in block 5
Hi! My Id is 0, I am the thread 0 out of 4 in block 0
Hi! My Id is 1, I am the thread 1 out of 4 in block 0
Hi! My Id is 2, I am the thread 2 out of 4 in block 0
Hi! My Id is 3, I am the thread 3 out of 4 in block 0
Hi! My Id is 16, I am the thread 0 out of 4 in block 4
Hi! My Id is 17, I am the thread 1 out of 4 in block 4
Hi! My Id is 18, I am the thread 2 out of 4 in block 4
Hi! My Id is 19, I am the thread 3 out of 4 in block 4
Hi! My Id is 12, I am the thread 0 out of 4 in block 3
Hi! My Id is 13, I am the thread 1 out of 4 in block 3
Hi! My Id is 14, I am the thread 2 out of 4 in block 3
Hi! My Id is 15, I am the thread 3 out of 4 in block 3
Hi! My Id is 4, I am the thread 0 out of 4 in block 1
Hi! My Id is 5, I am the thread 1 out of 4 in block 1
Hi! My Id is 6, I am the thread 2 out of 4 in block 1
Hi! My Id is 7, I am the thread 3 out of 4 in block 1

```

1.3 3 Memory Allocation

- Would it work the code without using the function “cudaMemcpy”.

Without cudaMemcpy calls, the device would not have access to the initial data in array a, and the modified data would not be transferred back to the host. As a result, the final print statement would not display the modified array a correctly.

- Change the value of “BLOCKSIZE” to, for instance, “4”. How does it affect the execution compared to the original output?

Original output was “Hello World!”. After setting BLOCKSIZE = 4 result is “Hello Worlo”. The program will modify characters in a by adding the matching integer values from array b. By this process, each adding will generate the given result, for instance $a[0] + b[0]$ will be ‘H’ + 15 = 72, 72 is the character ‘W’ in ASCII, ‘e’ + 10 = 111 which is ‘o’ and so on. However, the program will only be capable to read BLOCKSIZE elements of each and since BLOCKSIZE is set to 4, the threads within a block will have indices from 0 to 3, hence, the last characters from b will not be read and when added to a it will give us the remaining of Hello string, in this case o. BLOCKSIZE must have a minimum value of 6 in order to perform correctly.

```

[ ]: %%cu
#include <stdio.h>

const int N = 16;
const int GRIDSIZE = 1; //number of thread blocks
const int BLOCKSIZE = 4; //number of threads per thread block

__global__ void hello_decoder(char *a, int *b) {
    a[threadIdx.x] += b[threadIdx.x];
}

```

```

int main() {
    char a[N] = "Hello \0\0\0\0\0\0";
    int b[N] = {15, 10, 6, 0, -11, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0};
    char *ad;
    int *bd;
    const int csize = N*sizeof(char);
    const int isize = N*sizeof(int);

    printf("%s", a);

    cudaMalloc( (void**)&ad, csize );
    cudaMalloc( (void**)&bd, isize );
    cudaMemcpy( ad, a, csize, cudaMemcpyHostToDevice );
    cudaMemcpy( bd, b, isize, cudaMemcpyHostToDevice );

    hello_decoder<<<GRIDSIZE, BLOCKSIZE>>>(ad, bd);
    cudaMemcpy( a, ad, csize, cudaMemcpyDeviceToHost );
    cudaFree( ad );
    cudaFree( bd );

    printf("%s\n", a);
    return EXIT_SUCCESS;
}

```

Hello Worlo

1.4 4 Accelerator Characteristics

```
[ ]: ! nvidia-smi
```

Sun May 28 21:46:16 2023

```

+-----+
| NVIDIA-SMI 525.85.12      Driver Version: 525.85.12      CUDA Version: 12.0      |
+-----+-----+-----+-----+-----+
| GPU  Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                       |                    |    MIG M.   |
+-----+-----+-----+-----+-----+
|   0   Tesla T4              Off  | 00000000:00:04:0 Off  |            0         |
| N/A   39C    P8      9W / 70W | 0MiB / 15360MiB |      0%      Default  |
|                                       |                    |    N/A      |
+-----+-----+-----+-----+-----+

+-----+
| Processes:                                                       |
| GPU    GI    CI          PID    Type   Process name                      GPU Memory |

```

ID	ID	Usage
=====		
No running processes found		

Model	Architecture	SMs	Cores
Tesla T4	Turing (TU104)	40	2560

1.5 5 Programming Exercise

```
[ ]: %%writefile exercise.cu

#include <iostream>
#include <math.h>
#define VALUE 20
#define PROBLEMSIZE 1000000000

__global__ void add(float *x, float *y, int size) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < size)
        y[i] = x[i] + y[i];
}

int main(void) {
    float *x, *y;
    cudaMallocManaged(&x, PROBLEMSIZE*sizeof(float));
    cudaMallocManaged(&y, PROBLEMSIZE*sizeof(float));
    for (int i = 0; i < PROBLEMSIZE; i++) {
        float val = (float)(i % VALUE);
        x[i] = val;
        y[i] = (VALUE - val);
    }
    int blockSize = 256;
    int numBlocks = (PROBLEMSIZE + blockSize - 1) / blockSize;
    add<<<numBlocks, blockSize>>>(x, y, PROBLEMSIZE);
    cudaDeviceSynchronize();
    float error = 0.0f;
    for (int i = 0; i < PROBLEMSIZE; i++)
        error = fmax(error, fabs(y[i]-VALUE));
    if (error != 0)
        printf("Wrong result. Check your code, especially your kernel\n");
    cudaFree(x);
    cudaFree(y);
    return 0;
}
```

Writing exercise.cu

```
[ ]: ! nvcc -o exercise exercise.cu
```

```
[ ]: %%time  
! ./exercise
```

CPU times: user 122 ms, sys: 11.7 ms, total: 134 ms
Wall time: 18.9 s

- How different is managed transfers between CPU and GPU?

Managed transfers is handled automatically by the CUDA runtime system. This approach simplifies memory management by allowing the programmer to allocate and access memory using a unified memory address space. Unlike explicit memory transfers, where it has to be done manually data get copied between CPU and GPU memory, managed transfers automatically migrate data between CPU and GPU as needed. It also introduces some performance overhead due to the additional management and data migration operations. However, managed transfers provide convenience and flexibility by abstracting away low-level memory management tasks and enabling more efficient utilization of system resources in many scenarios and better programming experience.

- Check that it does not return an error (you can attach a screenshot).
- How long does it take to run (you can use the extension %%time at the beginning of the cell, or the Unix command time before the binary execution)?
- Provide block/thread configurations that reduces the execution time. Compare the configurations and the execution time (for instance, with a table).

```
[27]: %%writefile configurations.cu  
  
#include <iostream>  
#include <math.h>  
#include <stdio.h>  
#define VALUE 20  
#define PROBLEMSIZE 1000000000  
  
__global__ void add(float *x, float *y, int size) {  
    int i = threadIdx.x + blockIdx.x * blockDim.x;  
    if (i < size)  
        y[i] = x[i] + y[i];  
}  
  
int main(void) {  
    float *x, *y;  
    cudaMallocManaged(&x, PROBLEMSIZE * sizeof(float));  
    cudaMallocManaged(&y, PROBLEMSIZE * sizeof(float));  
    for (int i = 0; i < PROBLEMSIZE; i++) {  
        float val = (float)(i % VALUE);  
        x[i] = val;  
    }
```

```

    y[i] = (VALUE - val);
}

int configurations[][2] = {
    {256, 256},
    {512, 128},
    {1024, 64},
};

printf("Configuration\t| Execution Time (ms)\n");
printf("-----\n");

// Perform computations for each configuration and measure the execution time
for (int config = 0, numConfigurations = sizeof(configurations) /
    sizeof(configurations[0]); config < numConfigurations; config++) {
    int blockSize = configurations[config][0];
    int numBlocks = (PROBLEMSIZE + blockSize - 1) / blockSize;

    // Start measuring execution time
    cudaEvent_t start, end;
    cudaEventCreate(&start);
    cudaEventCreate(&end);
    cudaEventRecord(start);

    add<<<numBlocks, blockSize>>>(x, y, PROBLEMSIZE);
    cudaDeviceSynchronize();

    // Stop measuring execution time
    cudaEventRecord(end);
    cudaEventSynchronize(end);
    float milliseconds = 0;
    cudaEventElapsedTime(&milliseconds, start, end);

    // Print the configuration and its corresponding execution time
    printf("%d threads, %d blocks | %.2f\n", blockSize, numBlocks,
    milliseconds);
}

cudaFree(x);
cudaFree(y);
return 0;
}

```

Overwriting configurations.cu

```
[28]: ! nvcc -o configurations configurations.cu && ./configurations
```

```
Configuration | Execution Time (ms)
```

```
-----
256 threads, 3906250 blocks | 2034.24
512 threads, 1953125 blocks | 46.67
1024 threads, 976563 blocks | 46.72
```

- Is there any limitation on the blocks or threads quantities?

You can have up to $2^{31} - 1$ blocks in the x-dimension, and at most 65535 blocks in the y and z dimensions and a maximum of 1024 threads per block.

* [CUDA C++ Programming Guide - 6.2. Features and Technical Specifications](#)

1.6 6 Profiling

```
[ ]: ! nvprof ./exercise
```

```
==28754== NVPROF is profiling process 28754, command: ./exercise
==28754== Profiling application: ./exercise
==28754== Profiling result:
      Type  Time(%)      Time   Calls    Avg      Min      Max
Name
GPU activities: 100.00%  2.77920s        1  2.77920s  2.77920s  2.77920s
add(float*, float*, int)
  API calls:  77.21%  2.77921s        1  2.77921s  2.77921s  2.77921s
cudaDeviceSynchronize
      15.92%  572.99ms        2  286.50ms  236.12ms  336.87ms
cudaFree
      6.84%  246.17ms        2  123.09ms  51.124us  246.12ms
cudaMallocManaged
      0.02%  889.17us        1  889.17us  889.17us  889.17us
cuDeviceGetPCIBusId
      0.00%  113.31us       101  1.1210us    128ns  47.690us
cuDeviceGetAttribute
      0.00%  63.917us        1  63.917us  63.917us  63.917us
cudaLaunchKernel
      0.00%  23.675us        1  23.675us  23.675us  23.675us
cuDeviceGetName
      0.00%  1.5400us        3    513ns    160ns  1.1470us
cuDeviceGetCount
      0.00%    870ns        2    435ns    156ns    714ns
cuDeviceGet
      0.00%    535ns        1    535ns    535ns    535ns
cuModuleGetLoadingMode
      0.00%    492ns        1    492ns    492ns    492ns
cuDeviceTotalMem
      0.00%    257ns        1    257ns    257ns    257ns
cuDeviceGetUuid
```

==28754== Unified Memory profiling result:

Device "Tesla T4 (0)"

Count	Avg Size	Min Size	Max Size	Total Size	Total Time	Name
111732	67.225KB	4.0000KB	0.9961MB	7.163231GB	880.2939ms	Host To Device
22894	170.62KB	4.0000KB	0.9961MB	3.725292GB	341.6113ms	Device To Host
10846	-	-	-	-	2.758979s	Gpu page fault

groups

Total CPU Page faults: 34341

Provide the code of the non-managed CPU-GPU memory version of the problem.

```
[ ]: %%writefile non_managed.cu

#include <iostream>
#include <math.h>
#define VALUE 20
#define PROBLEMSIZE 1000

__global__ void add(float *x, float *y, int size) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < size)
        y[i] = x[i] + y[i];
}

int main(void) {
    float *x, *y;
    float *d_x, *d_y;
    int size = PROBLEMSIZE * sizeof(float);

    // Allocate memory on CPU
    x = (float*)malloc(size);
    y = (float*)malloc(size);

    // Initialize input arrays on CPU
    for (int i = 0; i < PROBLEMSIZE; i++) {
        float val = (float)(i % VALUE);
        x[i] = val;
        y[i] = (VALUE - val);
    }

    // Allocate memory on GPU
    cudaMalloc(&d_x, size);
    cudaMalloc(&d_y, size);

    // Copy input arrays from CPU to GPU
    cudaMemcpy(d_x, x, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_y, y, size, cudaMemcpyHostToDevice);

    // Launch kernel
```

```

int blockSize = 1;
int numBlocks = (PROBLEMSIZE + blockSize - 1) / blockSize;
add<<<numBlocks, blockSize>>>(d_x, d_y, PROBLEMSIZE);
cudaDeviceSynchronize();

// Copy result array from GPU to CPU
cudaMemcpy(y, d_y, size, cudaMemcpyDeviceToHost);

// Check for errors
float error = 0.0f;
for (int i = 0; i < PROBLEMSIZE; i++)
    error = fmax(error, fabs(y[i] - VALUE));
if (error != 0)
    printf("Wrong result. Check your code, especially your kernel\n");

// Free memory on GPU
cudaFree(d_x);
cudaFree(d_y);

// Free memory on CPU
free(x);
free(y);

return 0;
}

```

Overwriting non_managed.cu

```
[ ]: ! nvcc -o non_managed non_managed.cu && nvprof ./non_managed
```

==1322== NVPROF is profiling process 1322, command: ./non_managed

==1322== Profiling application: ./non_managed

==1322== Profiling result:

Name	Type	Time(%)	Time	Calls	Avg	Min	Max
GPU activities:		55.64%	7.4240us	1	7.4240us	7.4240us	7.4240us
add(float*, float*, int)		27.09%	3.6150us	2	1.8070us	1.6960us	1.9190us
[CUDA memcpy HtoD]		17.27%	2.3040us	1	2.3040us	2.3040us	2.3040us
[CUDA memcpy DtoH]		99.66%	340.34ms	2	170.17ms	8.2480us	340.33ms
API calls:		0.24%	804.43us	1	804.43us	804.43us	804.43us
cudaMalloc		0.04%	125.67us	101	1.2440us	142ns	52.676us
cuDeviceGetPCIBusId		0.03%	102.13us	2	51.063us	6.3910us	95.735us
cuDeviceGetAttribute							

cudaFree	0.02%	61.743us	3	20.581us	9.3790us	28.301us
cudaMemcpy	0.01%	38.888us	1	38.888us	38.888us	38.888us
cudaLaunchKernel	0.01%	25.743us	1	25.743us	25.743us	25.743us
cuDeviceGetName	0.00%	9.8090us	1	9.8090us	9.8090us	9.8090us
cudaDeviceSynchronize	0.00%	2.1110us	3	703ns	218ns	1.6630us
cuDeviceGetCount	0.00%	1.0790us	2	539ns	192ns	887ns
cuDeviceGet	0.00%	649ns	1	649ns	649ns	649ns
cuModuleGetLoadingMode	0.00%	512ns	1	512ns	512ns	512ns
cuDeviceTotalMem	0.00%	236ns	1	236ns	236ns	236ns
cuDeviceGetUuid						

- Which are the dominant API calls when the block size is equal to “1”? Why?

cudaMalloc with accounts for 99.66% of the API calls. The majority of the time is spent on allocating GPU memory because the block size. With a small block size, more individual memory allocations are required leading to that increased overhead.

- Would the dominant API calls be the same in the managed memory version? Why?

```
[ ]: ! nvcc -o managed exercise.cu && nvprof ./managed
```

```
==2678== NVPROF is profiling process 2678, command: ./managed
```

```
==2678== Profiling application: ./managed
```

```
==2678== Profiling result:
```

Name	Type	Time(%)	Time	Calls	Avg	Min	Max
GPU activities:	100.00%	3.56911s	1	3.56911s	3.56911s	3.56911s	3.56911s
add(float*, float*, int)	API calls:	81.99%	3.56914s	1	3.56914s	3.56914s	3.56914s
cudaDeviceSynchronize		12.17%	530.02ms	2	265.01ms	235.74ms	294.28ms
cudaFree		5.83%	254.02ms	2	127.01ms	58.169us	253.96ms
cudaMallocManaged		0.00%	118.48us	101	1.1730us	131ns	48.039us
cuDeviceGetAttribute		0.00%	74.836us	1	74.836us	74.836us	74.836us
cudaLaunchKernel		0.00%	25.015us	1	25.015us	25.015us	25.015us
cuDeviceGetName		0.00%	7.3720us	1	7.3720us	7.3720us	7.3720us

cuDeviceGetPCIBusId	0.00%	2.6080us	3	869ns	222ns	2.1200us
cuDeviceGetCount	0.00%	1.0220us	2	511ns	155ns	867ns
cuDeviceGet	0.00%	419ns	1	419ns	419ns	419ns
cuDeviceTotalMem	0.00%	408ns	1	408ns	408ns	408ns
cuModuleGetLoadingMode	0.00%	276ns	1	276ns	276ns	276ns
cuDeviceGetUuid						

==2678== Unified Memory profiling result:

Device "Tesla T4 (0)"

Count	Avg Size	Min Size	Max Size	Total Size	Total Time	Name
110850	67.672KB	4.0000KB	0.9961MB	7.154015GB	878.6278ms	Host To Device
22894	170.62KB	4.0000KB	0.9961MB	3.725292GB	343.8263ms	Device To Host
10926	-	-	-	-	3.550738s	Gpu page fault

groups

Total CPU Page faults: 34341

This time the dominants APIs are:

- cudaDeviceSynchronize (81.99%): the majority of the time of this code is spent waiting for the kernel to finish its execution since it is the primary computation in this code.
- cudaFree (12.17%) and cudaMallocManaged (5.83%) because large amount of memory is being allocated.